

# Day 2

## Prompt Engineering

- Focus on essential prompt engineering techniques
- Understand the function of the context window size
- Develop effective prompts for better model responses
- Explore advanced prompting strategies during Day 2



# What is Prompt Engineering?




- Crafting and refining inputs for large language models
- Involves structuring prompts to guide model behavior
- Focuses on achieving desired, high-quality model outputs
- Essential skill for maximizing the utility of LLMs

# Prompt Basics

## Definitions

- **Prompt:** An AI prompt is an **input or query** given to a large language model to elicit a specific response or output.
- **Prompt Engineering:** Prompt engineering is the practice of **designing and refining prompts** to optimize the responses generated by AI models



# Zero-Shot Prompting Explained

- Zero-shot prompting provides an instruction only
- The LLM uses its training to answer the request
- No specific examples are given within the prompt
- It relies solely on the model's existing knowledge



Classify the following review into neutral, negative or positive:

Review: "I absolutely loved this movie! The storyline was gripping and the acting was top-notch."

Sentiment: ...

# Zero-Shot Prompting Explained



# Few-Shot Prompting Explained



Classify the following movie reviews:

1. Review: "I absolutely loved this movie! The storyline was gripping and the acting was top-notch."  
Sentiment: Positive
2. Review: "The plot was boring and the characters were unconvincing."  
Sentiment: Negative
3. Review: "The movie was a bit slow, but the performances were excellent."

Sentiment: ...

- Providing examples guides the model's response
- Few-shot prompting includes several input-output pairs
- Improves accuracy and helps model learn specific tasks
- Examples are shown directly within the prompt's context

# Few-Shot Prompting Explained



# Zero-Shot vs. Few-Shot Prompting: A Comparative Guide

## Zero-Shot Prompting



### Best for...

- General knowledge tasks, creative writing, simple summaries.



### Speed

Faster to write (less typing).



### Cost (Tokens)

Cheaper (shorter prompts).




### Reliability

Moderate. The model might guess the format wrong.



## Few-Shot Prompting



### Best for...

- Specific formats, complex logic, mimicking a specific style/voice.



### Speed

Slower to write (need to curate examples).



### Cost (Tokens)

More expensive (examples use up token context).



### Reliability

High. The model follows the pattern provided.

Choosing the right prompting strategy depends on your task's complexity and the desired output precision.

# Building a Prompt Framework

## Prompt Basics

### Prompt components

A good prompt usually consists of:


- **Instruction:** A clear directive that specifies what the model should do.
- **Context:** Background/additional information that provides the necessary details to understand the task.
- **Input / question:** The specific query or data that the model needs to process.
- **Output type / format:** The desired structure or style of the response generated by the model.

Summarize the following article and write a list of top 3 important points in markdown format that answer the following question; “How is climate change affecting polar bear habitats and their ability to find food?”

**Article:**

<The article discusses the impact of climate change on polar bear populations in the Arctic.>

# Building a Prompt Framework




To perform Few-shot prompting effectively, you need a structure. Just throwing text at the model can confuse it. Use the **Standard Prompt Structure**:

1. **Role (Optional):** Who is the AI?
2. **Instruction:** What is the task?
3. **Examples (The "Shots"):** The pattern to follow.
4. **Context/Constraint:** Any guardrails?
5. **Input Data:** The actual thing to process.
6. **Output Indicator:** A cue for the AI to start.

# Context Window Deep Dive

- The context window is the model's short-term memory
- It defines the amount of text the model can process
- Window size limits the input and output length
- Exceeding the window causes the model to forget information



# What happens if the context window is exceeded ?



**Hands on .....**